



# SHOPPING LISTS ON THE CLOUD

*Large Scale Distributed Systems*

*Ana Beatriz Fontão - up202003574*

*José Luís Rodrigues - up202008462*

*Maria Sofia Gonçalves - up202006927*

# INTRODUCTION



Creation of a local-first shopping list application.



Data stored locally and on the Cloud.



Lists can be shared between clients.



CRUD operations over the lists' products.



Conflicts are resolved using CRDTs.



Scalable and highly-available system.

# INTRODUCTION

- What is being done?

An application that allows users to create shopping lists on their own devices, while at the same time collaborating on the cloud.

- How is it being done?

By using data versioning, sloppy quorum, consistent hashing, replication, a local-first and cloud-use approach and Conflict-free Replicated Data Types.

- Why is it being done?

To get the best of both worlds: seamless collaboration between users in different parts of the world (cloud) while maintaining full ownership of one's data (local-first).

## RELATED WORK

### Dynamo: Amazon's Highly Available Key-value Store.

- data versioning
- sloppy quorum
- consistent hashing
- replication

### Conflict-free Replicated Data Types


- data type whose replicas always converge, despite failures

### Local-first software

- offline functionalities as a primary concern
- incorporates cloud synchronization as a background task


# ARCHITECTURE

Single-Server




- + Easy to implement
- Limited to vertical scaling
  - SPF: low availability
- Vulnerable at increased loads

Peer-to-Peer



- + Highly scalable
- + Limited server interaction
- + Replication is done in the clients
  - Users are prone to disconnect (information gets lost)

Multi-server  
(our solution)



- + Vertically and Horizontally Scalable
- + High availability (clients can talk to all servers)
  - Loss in consistency (**SEC**)
  - Introduces complexity to ensure synchronization

# REPLICATION

- Usage of **Quorum Consensus** for reading and writing operations:
  - Lists are replicated at **N** nodes (servers).
  - Write operations are valid after **W** successful writes.
  - Read operations are valid after **R** successful writes.
- CAP Theorem:
  - Using Quorum increases the availability (A) and partition tolerance (P) of the system.
  - This comes at the cost of consistency (C).
  - Solution: Usage of CRDTs to guarantee **eventual consistency**.

# SHOPPING LIST

- The shopping list is composed by the elements:
  - **Id:** a 128-bit identifier for each list, ensuring that each Id is unique;
  - **Name:** the list's name, given by the user that created it;
  - **Products:** a listing of the desired products, explained in the following slides;
  - **Addwins:** composed by its elements, allows the merging of the list's products;
  - **Deleted:** a flag that is set to true if an user, who had access to the lists, decides to delete it.

```
{
  _id: ObjectId("6574e5fc16f19d4f130c0a18"),
  id: 'f3debc26-ffec-476a-bd58-02b6cbf0e0c4',
  name: 'Shop!',
  products: [
    {
      name: 'apple',
      pnCounter: {
        id: 'apple',
        inc: { id: 'apple', counters: { apple: 5 } },
        dec: { id: 'apple', counters: { apple: 2 } }
      },
      gCounter: { id: 'apple', counters: { apple: 2 } }
    },
    {
      name: 'potatos',
      pnCounter: {
        id: 'potatos',
        inc: { id: 'potatos', counters: { potatos: 7 } },
        dec: { id: 'potatos', counters: {} }
      },
      gCounter: { id: 'potatos', counters: {} }
    }
  ],
  addWins: {
    id: 'f3debc26-ffec-476a-bd58-02b6cbf0e0c4',
    cc: [
      { first: 'Sofia', second: Long("1") },
      { first: 'Luís', second: Long("3") },
      { first: 'Bea', second: Long("2") }
    ],
    set: [
      { first: 'Luís', second: 'potatos', third: Long("3") },
      { first: 'Sofia', second: 'apple', third: Long("1") }
    ],
    local_counter: Long("4")
  }
},
```

# CONSISTENCY (CRDTS)

- Context-free Replicated data Types (CRDTS) are a data type whose replicas always converge, despite any number of failures that may happen, thus ensuring strong eventual consistency.
- The following structures were used for the project:
  - Add-wins Set
  - PNCounter
  - GCounter



# ADD-WINS SET

- The Add-Wins Set is a CRDT that allows concurrent **addition of elements** ensuring consistency and avoiding conflicts.
- It contains three elements:
  - **Local counter:** A client private variable that increases on each operation
  - **Casual context:** Set of tuples that contain the client's id and their respective local counter. This set allows the system to keep track of the change history
  - **Set:** Set of triples that contain the client's id, a name of an element and their respective local counter. This set keeps track of the presence of elements
- **Merges** are performed based on the differences between the casual context and the set.

# G-COUNTER (GROW-ONLY)

- The G-Counter is a data structure that tracks the **incremental quantity** of an element.
- It contains two elements:
  - **Id**: Represents the id of the element.
  - **Counter**: Replica variable that can only increase and never decrease.
- The merges involves taking the maximum counter value out of all the replicas, ensuring that the total count is precise and never decreases over time.

# PN-COUNTER (POSITIVE-NEGATIVE)


- The PN-Counter is a CRDT similar to the G-Counter. The main difference is that this data structure allows quantities to be **increased and decreased**.
- Contains three elements:
  - **Id**: the product's name;
  - **Incremental counter**: a GCounter, responsible for informing about the quantity of the product that is added;
  - **Decremental counter**: a GCounter, responsible for informing about the quantity of the product that is removed;
- The merges are performed on each counter.
- The final quantity is given by the **difference** between the **incremental** counter and the **decremental** counter.

# CRDTS IMPLEMENTATION


- For this projects we combined all of the CRDT preciously presented.
- Add-Wins Set: This CRDT was used to keep track of the presence of a product in a shopping list.
- G-Counter: Used to represent the quantity of an item that is already bought (*quantityBought*). This number can only increase.
- PN-Counter: Indicates the quantity that should be bought (*quantity*) of a product in the shopping list. Can be decreased as well as increased.

# DATA PARTITIONING

- Used **Consistent Hashing** to distribute lists among servers
- All servers have full knowledge of the server grid
- Usage of a circular hash space with configurable size
- Node changes don't require a complete re-hashing
- Virtual Nodes are used to improve **load balancing**



# FAULT TOLERANCE




- Because the grid status is fully known, nodes continuously observe others status
- When a node doesn't receive an expected response for a certain interval, a fault is **detected**
- The node that first finds the fault is responsible to broadcast it
- **Hinted Handoff:** After a status change, nodes relocate lists that belong to the sick node, to ensure **replication**
- The image shows the failure of node A and subsequent relocation of List1 to the next available physical node (C)

# NODE RECOVERY

- Upon receiving a hinted handoff, the node keeps track which lists are only temporarily relocated
- When a sick node comes back online, all hinted lists are returned, and changes are updated

For **new nodes** joining the network, the process is analogous:

- The new node sends a heartbeat to all the other nodes
- Upon detecting the new node, the servers send it the lists that now belong to it



# IMPLEMENTATION

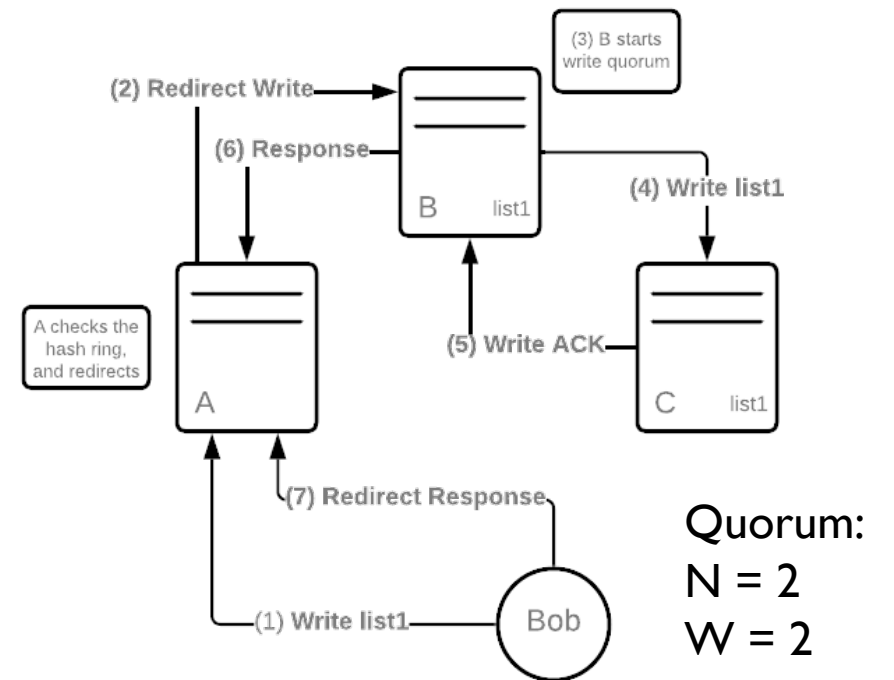
- The program was built in **Java**
- Communication is done with ZMQ sockets
- There are two types of processes: **nodes** (servers) and **clients**
- Client-Server communication is done DEALER-ROUTER
  - Client is able to fairly distribute requests among known nodes
  - Servers can talk to multiple clients at the same time
- Server-Server communication is done ROUTER-ROUTER
  - This allows for a star topology and fully interconnected nodes



# PROGRAM EXECUTION

## LIST WRITE OPERATION

1. Client sends request in a **round-robin** fashion to each server
2. When the server receives a request, this may be redirected, when the node is out of the ring
3. The operation begins by starting a quorum.
4. Node B tells C to write the list
5. Upon receiving confirmation from C, quorum is achieved
6. The response is redirected to A
7. A now sends the reply to the client



# PARAMETER TUNING

The program is built to allow the configuration of parameters, including:

- **Number of virtual nodes**
  - Increasing the number of virtual nodes improves load balancing, by increasing the granularity in distributing data across the hash ring.
  - May have diminishing returns as the hashing complexity increases
- **Ring Size**
  - Ring size can be changed to better suit each application's dimensionality.
  - Increasing this value diminishes the number of collisions but hashes become more sparse
- **Quorum**
  - Varying quorum values  $N$ ,  $W$  and  $R$  allows the developer to empower CAP properties best suit to a specific application
  - Higher  $N$  value enhances partition tolerance by increasing the number of replicas
  - Increasing  $R$  improves consistent at the cost of availability
- **Timeouts**
  - Configuring delays changes how fast the system reacts to faults

# DISCUSSION

- Using a static grid configuration simplifies communications but discourages online modifications. Mechanisms like **dynamic discovery** or the usage of **brokers** can help in solving this problem.
- The system relies on each node having complete network knowledge. This can, however, impose a limit on horizontal scaling as the overhead in maintaining this state increases.
- The focus on availability and the local-first structure gives the user an always ready to use application.
  - The trade-off is a considerable lost in consistency.
  - Using **eventual consistency** reduces user experience in the sense that changes may not be updated as fast as in more online-oriented architectures

# CONCLUSION

- **Scalability:** The product was build with a focus on scalability and allows the addition of multiple servers to the network.
- **Local-First:** An offline-oriented architecture was achieved by prioritizing local actions and updating with the cloud on the background.
- **Collaboration:** By writing operations to the cloud, users are given the option to work on the same data at the same time.
- **Consistency:** The usage of CRDTs empowered the program with the ability to merge data and ensure seamless colaboration.